

MACHINE LEARNING · FINAL PROJECT

Verified real-work screening for engineering hiring in the AI era

Real work. Real person. Provable. One explainable, audit-ready score - with AI use allowed, not policed.

Ayush Jain · Christoph Dittrich · Mark Garvin Jeyaraj | Northeastern University

THE PROBLEM

Generative AI broke the coding screen

- Puzzle interviews are solved instantly by an LLM - a passing score no longer signals the **candidate's** ability.
- "AI-detection" is an unwinnable arms race: high false positives punish honest candidates and destroy trust.
- Pipelines flood with indistinguishable applicants; scarce **senior-engineer time** is wasted on the wrong people.

Our bet: stop detecting AI. Measure real, role-specific work *with AI allowed*, and turn it into a reliable, explainable signal.

RELATED WORK

What exists, and where we differ

Industry

- Auto-graded platforms (HackerRank, Codility) - the puzzle screens AI undermined.
- AI-proctoring / detection vendors - surveillance + brittle classifiers.

Research we build on

- **LLM-as-a-judge** (Zheng et al., MT-Bench, NeurIPS 2023): judges agree with humans about as often as humans agree with each other.
- **HumanEval** (Chen et al., 2021): execution as objective ground truth for code.

We take the LLM-judge idea but **remove the score from the model's control** and ground correctness in execution.

Inputs, and how we test without real outcomes

- **Live inputs:** a messy work sample (code + rationale) · a versioned rubric · server-side behavioral signals (typed-vs-pasted, edit timing, tab-focus) · recruiter-reported outcomes (advanced/offer/hired/retained). No biometrics.
- **Scorer-validity corpus:** 200 execution-labeled programs from HumanEval - 100 correct, 100 systematically mutated (boundary flips, off-by-one, subtle LLM bugs).
- **The cold-start problem:** outcome labels are exactly what a new customer lacks. So we drive the two evaluations with **seeded Monte-Carlo simulation** - reproducible, and it sweeps the data-scarcity axis directly.


METHODS

Hybrid scorer + two statistical guarantees

- **Scoring (LLM-as-judge, code tally):** the model emits only per-criterion points + evidence; **our code** computes the 0-100 score. Injection can't move arithmetic. K=3 self-consistency; correctness grounded in execution; every decision hash-chained.
- **Cold-start calibration:** a Beta-Binomial (empirical-Bayes) posterior shrinks a band's rate toward a documented prior; tiered **prior** **provisional** **calibrated** with a 90% credible interval that narrows as data arrives.
- **Subgroup fairness:** EEOC four-fifths (80%) rule on opt-in labels - run on synthetic cohorts now, before real data.

RESULTS · 1 OF 2

Calibration that works from day one



Shrinkage vs. raw rate as labels accrue (4,000 trials/band).

-56.6%

error when labels are scarce
(MAE 20.1 → 8.7 pp)

100%


available vs. 0% for the raw band
below the floor

95.1%

90% credible-interval coverage
(honest uncertainty)

RESULTS · 2 OF 2

Subgroup fairness, checked early




Impact ratio by subgroup; 0.80 four-fifths threshold.

- **Blind scorer passes:** every subgroup ≥ 0.80 (min 0.87).
- **Injected bias is caught:** penalized group flagged at 0.35.
- **Privacy floor holds:** a 3-person group is suppressed, not exposed.

DEMO

The working application



Live at **/demo/features**: the real calibration badge sharpening from prior → provisional → calibrated, and the fairness panel (blind passes, biased flagged).

Take-aways

- **One decision drove everything:** never let the model produce the number. Auditability, reproducibility, injection-resistance, and testability all follow.
- **"Not enough data" is a quantity, not a switch:** report a shrunk estimate with an interval, and fairness can be checked before the first real candidate.
- **Next:** hierarchical priors that pool across roles; scorer-human agreement on a public corpus; live four-fifths alerting.

710 tests green. Code: github.com/ayushjain-10/touchstone